

# yz709-ds-sup1

[Question 1](#)

[Question 2](#)

[Question 4](#)

[Question 5](#)

[Question 6](#)

[Question 7](#)

[Question 8](#)

## Question 1

**Question 1.** Given a dataset  $(x_1, \dots, x_n)$ , we wish to fit a Poisson distribution. This is a discrete random variable with a single parameter  $\lambda > 0$ , called the rate, and

$$\Pr(x; \lambda) = \frac{\lambda^x e^{-\lambda}}{x!} \quad \text{for } x \in \{0, 1, 2, \dots\}.$$

Show that the maximum likelihood estimator for  $\lambda$  is  $\hat{\lambda} = n^{-1} \sum_{i=1}^n x_i$ .

$$\begin{aligned} (x_1, \dots, x_n) \\ \Pr(x; \lambda) &= \frac{\lambda^x e^{-\lambda}}{x!} \quad x \in \{0, 1, 2, \dots\} \\ \log \Pr(x; \lambda) &= x \log(\lambda) - \lambda - \log(x!) \\ \log \text{lik} &= \sum_{i=1}^n \log \Pr(x_i; \lambda) = \sum_{i=1}^n x_i (\log \lambda) - \lambda n - \sum_{i=1}^n \log(x_i!) \\ \text{when } \frac{\partial \log \text{lik}}{\partial \lambda} &= \frac{\sum_{i=1}^n x_i}{\lambda} - n = 0 \\ \Rightarrow \hat{\lambda} &= n^{-1} \sum_{i=1}^n x_i \end{aligned}$$



Comments:

## Question 2

**Question 2.** Given a dataset [3,2,8,1,5,0,8], we wish to fit a Poisson distribution. Give code to achieve this fit, using `scipy.optimize.fmin`.

$$X \sim \text{Poi}(\lambda), \lambda > 0$$
$$P(X=x) = \frac{\lambda^x e^{-\lambda}}{x!}$$

since `scipy.optimize.fmin` needs parameters  $\in \mathbb{R}$   
let  $\lambda = e^t, t \in \mathbb{R}$

```
x = np.array(3, 2, 8, 1, 5, 0, 8)

def loglik(x, t):
    return np.log(scipy.stats.poisson.pmf(x, t))

t_hat = scipy.optimize.fmin(lambda l: -np.sum(loglik(x, np.log(l))), np.log(0.2))
l_hat = np.exp(t_hat) # l_hat is the optimised lambda value
```



Comments:

## Question 4

**Question 4.** Given a dataset  $(x_1, \dots, x_n)$ , we wish to fit the  $\text{Uniform}[0, \theta]$  distribution, where  $\theta$  is unknown. Show that the maximum likelihood estimator is  $\hat{\theta} = \max_i x_i$ .

$$P(X \leq x) = \begin{cases} 0 & , x < 0 \\ \frac{x}{\theta} & , x \in [0, \theta] \\ 1 & , x > \theta \end{cases}$$

$$P_{\Gamma}(x|\theta) = \frac{d}{dx} P(X \leq x) = \begin{cases} 0 & , x < 0 \\ \frac{1}{\theta} & , x \in [0, \theta] \\ 0 & , x > \theta \end{cases}$$

$$= \frac{1}{\theta} \mathbb{I}_{x \geq 0} \mathbb{I}_{\theta \geq x}$$

$$\prod_{i=1}^n P_{\Gamma}(x_i|\theta) = \frac{1}{\theta^n} \prod_{i=1}^n \mathbb{I}_{x_i \geq 0} \mathbb{I}_{\theta \geq x_i}$$

hence at  $\max_i x_i$ , all indicator functions would be 1, hence  $\prod_{i=1}^n P_{\Gamma}(x_i|\theta) = \frac{1}{\theta^n}$ , where  $\theta \geq \max_i x_i$

since  $\frac{1}{\theta^n}$  decreases as  $\theta$  increases, we select the smallest  $\theta \in [\max_i x_i, \infty)$   
 $\Rightarrow \hat{\theta} = \max_i x_i$



Comments:

## Question 5

**Question 5 (A/B testing).** Your company has two systems which it wishes to compare,  $A$  and  $B$ . It has asked you to compare the two, on the basis of performance measurements  $(x_1, \dots, x_m)$  from system  $A$  and  $(y_1, \dots, y_n)$  from system  $B$ . Any fool using Excel can just compare the averages,  $\bar{x} = m^{-1} \sum_{i=1}^m x_i$  and  $\bar{y} = n^{-1} \sum_{i=1}^n y_i$ , but you are cleverer than that and you will harness the power of Machine Learning.

Suppose the  $x_i$  are drawn from  $X \sim \text{Normal}(\mu, \sigma^2)$ , and the  $y_i$  are drawn from  $Y \sim \text{Normal}(\mu + \delta, \sigma^2)$ , and all the samples are independent, and  $\mu$ ,  $\delta$ , and  $\sigma$  are unknown. Find maximum likelihood estimators for the three unknown parameters.

Q5

$$Pr(x_i | \mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x_i - \mu)^2}{2\sigma^2}}$$

$$Pr(y_i | \mu, \delta, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(y_i - (\mu + \delta))^2}{2\sigma^2}}$$

$$\text{loglik}_x = \sum_{i=1}^m \log Pr(x_i | \mu, \sigma) = -\frac{m}{2} \log(2\pi\sigma^2) - \sum_{i=1}^m \frac{(x_i - \mu)^2}{2\sigma^2}$$

$$\text{loglik}_y = \sum_{i=1}^n \log Pr(y_i | \mu, \delta, \sigma) = -\frac{n}{2} \log(2\pi\sigma^2) - \sum_{i=1}^n \frac{(y_i - (\mu + \delta))^2}{2\sigma^2}$$

$$\frac{\partial \text{loglik}_x}{\partial \mu} = \sum_{i=1}^m \frac{x_i - \mu}{\sigma^2} = 0 \quad \frac{\partial \text{loglik}_x}{\partial \sigma} = -\frac{m}{2\pi\sigma} + \sum_{i=1}^m \frac{x_i^2}{\sigma^3} = 0$$

$$\Rightarrow \mu = \frac{\sum_{i=1}^m x_i}{m} = \bar{x}$$

$$\sigma^2 = \frac{2\pi}{m^2} \sum_{i=1}^m (x_i - \bar{x})^2$$

$$\frac{\partial \text{loglik}_y}{\partial \mu} = \sum_{i=1}^n \frac{y_i - \mu - \delta}{\sigma^2} = 0$$

$$\Rightarrow \sigma = \frac{\sqrt{m}}{m} \sqrt{\sum_{i=1}^m (x_i - \bar{x})^2}$$

$$\Rightarrow \mu = \frac{\sum_{i=1}^n y_i - \delta}{n} = \bar{y} - \frac{\delta}{n}$$

$$\Rightarrow \delta = \sum_{i=1}^n y_i - n\bar{x}$$



Comments:

## Question 6

**Question 6.** Let  $x_i$  be the population of city  $i$ , and let  $y_i$  be the number of crimes reported. Consider the model  $Y_i \sim \text{Poisson}(\lambda x_i)$ , where  $\lambda > 0$  is an unknown parameter. Find the maximum likelihood estimator  $\hat{\lambda}$ .

$$Y_i \sim \text{Poisson}(\lambda x_i), \lambda > 0$$

$$Pr(Y_i = y_i | \lambda x_i) = \frac{(\lambda x_i)^{y_i} e^{-\lambda x_i}}{y_i!}$$

$$\log Pr(y_i | \lambda x_i) = y_i \log(\lambda x_i) - \lambda x_i - \log(y_i!)$$

$$\log \text{lik} = \sum_{j=1}^n \log Pr(y_j | \lambda x_j) = [\log(\lambda) + \log(x_j)] \sum_{j=1}^n y_j - \lambda \sum_{j=1}^n x_j - \sum_{j=1}^n \log(y_j!)$$

$$\frac{\partial \log \text{lik}}{\partial \lambda} = \frac{\sum_{j=1}^n y_j}{\lambda} - \sum_{j=1}^n x_j = 0$$

$$\hat{\lambda} = \bar{x} \bar{y}$$


population or city  $i$   $\leftrightarrow$  the average number of crimes



Comments:

## Question 7

**Question 7.** We wish to fit a piecewise linear line to a dataset, as shown below. The inflection point is given, and we wish to estimate the slopes and intercepts. Explain how to achieve this using a linear modelling approach.




Note. As a sanity check, you should implement your model formula as a function and plot it. Here's a function that **fails** the check.

```
def pred(x, m1, c1, m2, c2, inflection_x=3):
    e = numpy.where(x <= inflection_x, 1, 0)
    return e*(m1*x + c1) + (1-e)*(m2*x+c2)
x = numpy.linspace(0,5,1000)
plt.plot(x, pred(x, m1=0.5, c1=0, m2=1, c2=2))
```

$$y = \beta_0 + \beta_1 x + \beta_2 (x - 3) 1_{x \geq 3}$$

- When  $x < 3$ , we have  $y = \beta_0 + \beta_1 x$  because  $1_{x \geq 3} = 0$
- When  $x \geq 3$ , we have  $y = \beta_0 + \beta_1 x + \beta_2 (x - 3) = (\beta_0 - 3\beta_2) + (\beta_1 + \beta_2)x$



```
import scipy, seaborn as sns
def f(β, x, y, knot):
    β0, β1, β2 = β
    return np.sum(np.power(y - (β0 + β1*x + β2*(x-knot)*np.where(x>=3, 1, 0)), 2))
def pred(β, x, knot):
    β0, β1, β2 = β
    return β0 + β1*x + β2*(x-knot)*np.where(x>=3, 1, 0)

x = np.linspace(0,5,1000)
y = np.array([np.where(i>=3, 3 * i + 1 + np.random.normal(0, 1), 2 * i - 4 + np.random.normal(0, 1)) for i in x])
β = scipy.optimize.fmin(lambda β: f(β, x, y, 3), (1,1,1), maxiter=100)

sns.scatterplot(x,y,size=3)
sns.lineplot(x, pred(β,x,3))
plt.show()
```



Comments:

## Question 8

**Question 8.** For the climate data from section 2.2.5 of lecture notes, we proposed the model

$$\text{temp} \approx \alpha + \beta_1 \sin(2\pi t) + \beta_2 \cos(2\pi t) + \gamma t$$

in which the  $+\gamma t$  term asserts that temperatures are increasing at a constant rate. We might suspect though that temperatures are increasing non-linearly. To test this, we can create a non-numerical feature out of  $t$  by

$$u = \text{'decade\_'} + \text{str}(\text{math.floor}(t/10)) + \text{'0s'}$$

(which gives us values like 'decade\_1980s', 'decade\_1990s', etc.) and fit the model

$$\text{temp} \approx \alpha + \beta_1 \sin(2\pi t) + \beta_2 \cos(2\pi t) + \gamma_u.$$

Write this as a linear model, and give code to fit it. *[Note. You should explain what your feature vectors are, then give a one-line command to estimate the parameters.]*


- Set up a base line model with no Secular Trend:  $y = \alpha + \beta_1 * \sin(2\pi t) + \beta_2 * \sin(2\pi t)$
- Assume we add on a variable  $\beta_3$  which changes every 10 years
- Then split the years into [1959-1968], [1969-1978], [1979-1988], [1989-1998], [1999-2008], [2009-2018], [2019-2021], each with a linear model  $y = \alpha + \beta_1 * \sin(2\pi t) + \beta_2 * \sin(2\pi t) + \beta_3 = (\alpha + \beta_3) + \beta_1 * \sin(2\pi t) + \beta_2 * \sin(2\pi t)$ , where  $\beta_3$  is different and all other coefficients are the same and the feature vectors are `[np.ones(len(t), np.sin(2*np.pi*t), np.cos(2*np.pi*t))]`

```
# a base linear model which uses all data from year 1959 to year 2021
# y = alpha + beta1 * sin(2*pi*t) + beta2 * sin(2*pi*t)
t,y = data[0]
X = np.column_stack([np.sin(2*np.pi*t), np.cos(2*np.pi*t)])
model = sklearn.linear_model.LinearRegression()
model.fit(X, y)
coef_base, (beta1,beta2) = (model.intercept_, model.coef_)

coef = []
pred = []

# minimising error term squared
def f(v, x_val, y_val):
    return np.sum(np.power(y_val - (v + beta1 * np.sin(2*np.pi*x_val) + beta2 * np.cos(2*np.pi*x_val)),2));

for x,y in data:
    new_alpha = scipy.optimize.fmin(lambda v : f(v, x, y), coef_base, maxiter=100); # find the optimal coefficient alpha
    coef.append((new_alpha[0],beta1,beta2))
    t = np.linspace(np.min(x), np.max(x), 1000)
    Xnew = np.column_stack([np.sin(2*np.pi*t), np.cos(2*np.pi*t)])
    t_pred = new_alpha + beta1 * np.sin(2*np.pi*t) + beta2 * np.cos(2*np.pi*t);
    pred.append((t, t_pred))
```



```
# output:
"""
from year 1959 to year 1968, coef difference to the base line model (β3): 0.000
from year 1969 to year 1978, coef difference to the base line model (β3): 0.162
from year 1979 to year 1988, coef difference to the base line model (β3): -0.110
from year 1989 to year 1998, coef difference to the base line model (β3): 0.781
from year 1999 to year 2008, coef difference to the base line model (β3): 1.271
from year 2009 to year 2018, coef difference to the base line model (β3): 1.070
from year 2019 to year 2021, coef difference to the base line model (β3): 1.326
"""
```



Comments: